

Yaniv Sheena

308446764

- 1) I generated training and test sets, each consist of 10,000 examples.
- 2) The average length of the sequences is ~140, where the size of each part of the sequence (i.e. [1-9] sequences and letter sequences) is randomly chose in the range: [1,30].
- 3) My network managed to distinguish the two languages (I.e. it reached 100% accuracy on the test-set) after one iteration, that took 58 seconds (I'm running on a VM with relatively low resources).
- 4) In order to make my network work:
 - a. I used a 1-layer LSTM with input of 50 dimensions and output of 50 dimensions as well.
 - b. On top of the LSTM, I built a MLP with 1 hidden layer of 20 dimensions (using tanh as activation function), and output layer consists 2 outputs (binary classifier).
 - c. I used learning-rate = 1